

titanic-data-analysis

May 11, 2023

```
[2]: # importing the required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import math
```

```
[13]: import warnings
warnings.filterwarnings("ignore")
```

```
[4]: titanic = pd.read_csv("titanic.csv")
```

```
[5]: titanic.head()
```

```
[5]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
[6]: # number of passengers travelling in the ship
print("no of passengers are --->>>",len(titanic))
```

no of passengers are --->>> 891

```
[7]: titanic.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Sex             891 non-null   object
5   Age            714 non-null   float64
6   SibSp          891 non-null   int64
7   Parch          891 non-null   int64
8   Ticket         891 non-null   object
9   Fare           891 non-null   float64
10  Cabin          204 non-null   object
11  Embarked       889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
[9]: # check for duplicate data
titanic.duplicated().sum()    # no duplicate data
```

[9]: 0

```
[10]: # check for missing values
titanic.isnull().sum()
```

```
[10]: PassengerId      0
Survived             0
Pclass               0
Name                 0
Sex                  0
Age                 177
SibSp                0
Parch                0
Ticket               0
Fare                 0
Cabin                687
Embarked             2
```


dtype: int64

```
[11]: # Age, Cabin and Embarked columns have missing values
# checking the % of missing data
titanic.isnull().mean()*100
```

```
[11]: PassengerId    0.000000
Survived         0.000000
Pclass           0.000000
Name             0.000000
Sex              0.000000
Age             19.865320
SibSp            0.000000
Parch           0.000000
Ticket           0.000000
Fare             0.000000
Cabin           77.104377
Embarked         0.224467
dtype: float64
```

```
[12]: # dropping the Cabin column because it has more than 75% of mv
titanic.drop("Cabin", axis=1, inplace=True)
```

```
[15]: # imputing missing values in age column
sns.boxplot(titanic.Age)
plt.show()
# the boxplot is showing more than 65 values are outliers but the age
# should be near 80 so it should be ok
```



```
[16]: titanic.Age.mean()
```

```
[16]: 29.69911764705882
```

```
[17]: # imputing the Age nan values with mean  
titanic.Age = titanic.Age.replace({np.nan:30})
```

```
[18]: titanic.Age.head() # it is in float so converting to integer
```

```
[18]: 0    22.0  
     1    38.0  
     2    26.0  
     3    35.0  
     4    35.0  
     Name: Age, dtype: float64
```

```
[19]: titanic.Age.isnull().sum()
```

```
[19]: 0
```

```
[22]: titanic.Age = titanic.Age.astype(np.int64)
```

```
[23]: titanic.Age.dtype
```

```
[23]: dtype('int64')
```

```
[25]: # Imputing the missing values of Embarked column
# the column is object so replacing nan values with mode
titanic.Embarked.value_counts()
```

```
[25]: S    644
      C    168
      Q     77
      Name: Embarked, dtype: int64
```

```
[26]: titanic.Embarked = titanic.Embarked.replace({np.nan:"S"})
```

```
[27]: titanic.isnull().sum()    # no missing values
```

```
[27]: PassengerId    0
      Survived      0
      Pclass       0
      Name         0
      Sex          0
      Age          0
      SibSp        0
      Parch        0
      Ticket       0
      Fare         0
      Embarked     0
      dtype: int64
```

```
[28]: numeric = []
      categ = []
      for i in titanic.columns:
          if titanic[i].dtype == "int64" or titanic[i].dtype == "float64":
              numeric.append(i)
          else:
              categ.append(i)
```


```
[29]: numeric
```

```
[29]: ['PassengerId', 'Survived', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare']
```


```
[30]: categ
```

```
[30]: ['Name', 'Sex', 'Ticket', 'Embarked']
```


```
[39]: sns.countplot(titanic.Sex, hue=titanic.Survived)
      plt.show() # more females are survived compare to male
```




```
[40]: sns.countplot(titanic.Embarked, hue=titanic.Survived)  
plt.show() # most of the passengers are ported from S
```




```
[41]: sns.countplot(titanic.Pclass, hue=titanic.Survived)
plt.show() # passenger of class 1 are survived more than other class
# passengers who are not survived are majorly from class 3
```




```
[43]: titanic.Age.plot.hist()
plt.show() # most of the passengers are having age between 20 to 40
```




```
[45]: sns.countplot(titanic.Survived, hue=titanic.Sex)
plt.show() # out of 891 ,around 340 are survived
# in surviving females get more weightage than males
```




```
[46]: sns.countplot(titanic.Parch)
plt.show() # most of them are without any children or parents
```

```
[47]: sns.countplot(titanic.SibSp)
plt.show() # most of them are without any siblings or spouse
```



```
[49]: sns.heatmap(titanic.corr(), annot=True, cmap="Greens")
plt.show()
```



```
[50]: titanic.columns
```

```
[50]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
          'Parch', 'Ticket', 'Fare', 'Embarked'],
          dtype='object')
```

```
[52]: # create dummies
# column "Pclass" it has 3 values 1,2,3
pcl = pd.get_dummies(titanic.Pclass, drop_first=True)
pcl.head()
```

```
[52]:    2  3
0  0  1
1  0  0
2  0  1
3  0  0
4  0  1
```

```
[53]: # sex column
# replacing values of males with 1 and female with 0
titanic.Sex = titanic.Sex.replace({"male":1, "female":0})
```

```
[54]: # Embarked column
embark = pd.get_dummies(titanic.Embarked, drop_first=True)
embark.head()
```

```
[54]:      Q  S
0  0  1
1  0  0
2  0  1
3  0  1
4  0  1
```

```
[55]: titanic_new = pd.concat([titanic,pcl,embark], axis=1)
```

```
[56]: # dropping the useless columns
titanic_new.drop(['PassengerId', 'Pclass', 'Name', 'Ticket', 'Embarked'],axis=1,
                 inplace=True)
```

```
[57]: titanic_new.head()
```

```
[57]:      Survived  Sex  Age  SibSp  Parch      Fare    2    3    Q    S
0           0    1   22     1      0    7.2500    0    1    0    1
1           1    0   38     1      0   71.2833    0    0    0    0
2           1    0   26     0      0    7.9250    0    1    0    1
3           1    0   35     1      0   53.1000    0    0    0    1
4           0    1   35     0      0    8.0500    0    1    0    1
```

Train and Test data

```
[58]: X = titanic_new.drop("Survived",axis=1)
y = titanic_new.Survived
```

```
[61]: from sklearn.model_selection import train_test_split
```

```
[136]: x_train,x_test,y_train,y_test = train_test_split(X,y, train_size=0.75,
↳random_state=100)
```

```
[137]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
[137]: ((668, 9), (223, 9), (668,), (223,))
```

```
[138]: from sklearn.linear_model import LogisticRegression
```

```
[139]: lr = LogisticRegression()
```

```
[140]: lr.fit(x_train, y_train)
```

```
[140]: LogisticRegression()
```

```
[142]: y_pred = lr.predict(x_test)
```

```
[143]: y_pred
```

```
[143]: array([1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1,
         0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0,
         0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
         0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1,
         0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0,
         0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0,
         0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1,
         0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1,
         0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0,
         1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0,
         0, 1, 1], dtype=int64)
```

```
[145]: lr.predict_proba(x_test)
```

```
[145]: array([[0.24498507, 0.75501493],
         [0.29740863, 0.70259137],
         [0.90555105, 0.09444895],
         [0.29199896, 0.70800104],
         [0.79719063, 0.20280937],
         [0.25138113, 0.74861887],
         [0.91273187, 0.08726813],
         [0.79429422, 0.20570578],
         [0.07201488, 0.92798512],
         [0.66023096, 0.33976904],
         [0.02991357, 0.97008643],
         [0.88977604, 0.11022396],
         [0.80787651, 0.19212349],
         [0.58620948, 0.41379052],
         [0.4745412 , 0.5254588 ],
         [0.84251599, 0.15748401],
         [0.09256795, 0.90743205],
         [0.66826481, 0.33173519],
         [0.60627472, 0.39372528],
         [0.88592897, 0.11407103],
         [0.91466345, 0.08533655],
         [0.05060341, 0.94939659],
         [0.87834135, 0.12165865],
         [0.3604996 , 0.6395004 ],
         [0.9082167 , 0.0917833 ]],
```

[0.9146019 , 0.0853981],
[0.24124133, 0.75875867],
[0.90540386, 0.09459614],
[0.92794392, 0.07205608],
[0.85909046, 0.14090954],
[0.99116047, 0.00883953],
[0.1930689 , 0.8069311],
[0.60161681, 0.39838319],
[0.74004035, 0.25995965],
[0.07503033, 0.92496967],
[0.92528054, 0.07471946],
[0.10474279, 0.89525721],
[0.08034538, 0.91965462],
[0.25478045, 0.74521955],
[0.91471164, 0.08528836],
[0.41711692, 0.58288308],
[0.04651352, 0.95348648],
[0.84453918, 0.15546082],
[0.87182948, 0.12817052],
[0.99116047, 0.00883953],
[0.33230373, 0.66769627],
[0.8858622 , 0.1141378],
[0.04336588, 0.95663412],
[0.82028211, 0.17971789],
[0.14463011, 0.85536989],
[0.87011604, 0.12988396],
[0.72550251, 0.27449749],
[0.92322973, 0.07677027],
[0.9027845 , 0.0972155],
[0.91763461, 0.08236539],
[0.87100174, 0.12899826],
[0.92515657, 0.07484343],
[0.92166276, 0.07833724],
[0.82999979, 0.17000021],
[0.78005803, 0.21994197],
[0.78143661, 0.21856339],
[0.64673872, 0.35326128],
[0.88937431, 0.11062569],
[0.8999849 , 0.1000151],
[0.74004035, 0.25995965],
[0.35181033, 0.64818967],
[0.88589879, 0.11410121],
[0.23133112, 0.76866888],
[0.18598234, 0.81401766],
[0.93015661, 0.06984339],
[0.7735693 , 0.2264307],
[0.94104645, 0.05895355],

[0.9200851 , 0.0799149],
[0.38588116, 0.61411884],
[0.90894035, 0.09105965],
[0.79959657, 0.20040343],
[0.65013505, 0.34986495],
[0.87694862, 0.12305138],
[0.93198589, 0.06801411],
[0.87695094, 0.12304906],
[0.74249147, 0.25750853],
[0.81402474, 0.18597526],
[0.16223886, 0.83776114],
[0.11709219, 0.88290781],
[0.15218735, 0.84781265],
[0.89989021, 0.10010979],
[0.9111617 , 0.0888383],
[0.15638383, 0.84361617],
[0.70410623, 0.29589377],
[0.95787246, 0.04212754],
[0.89316678, 0.10683322],
[0.57323484, 0.42676516],
[0.91466345, 0.08533655],
[0.88945184, 0.11054816],
[0.05109101, 0.94890899],
[0.39831677, 0.60168323],
[0.26818811, 0.73181189],
[0.09717053, 0.90282947],
[0.9146019 , 0.0853981],
[0.34251262, 0.65748738],
[0.89997915, 0.10002085],
[0.89470124, 0.10529876],
[0.88179455, 0.11820545],
[0.08339893, 0.91660107],
[0.9200851 , 0.0799149],
[0.17809994, 0.82190006],
[0.92760136, 0.07239864],
[0.38588116, 0.61411884],
[0.32786604, 0.67213396],
[0.69598802, 0.30401198],
[0.93002365, 0.06997635],
[0.83632069, 0.16367931],
[0.47785329, 0.52214671],
[0.89316678, 0.10683322],
[0.39819454, 0.60180546],
[0.89294829, 0.10705171],
[0.9305139 , 0.0694861],
[0.95187078, 0.04812922],
[0.89997915, 0.10002085],

[0.43016626, 0.56983374],
[0.90588574, 0.09411426],
[0.9146019 , 0.0853981],
[0.11235038, 0.88764962],
[0.72776156, 0.27223844],
[0.97419364, 0.02580636],
[0.12804759, 0.87195241],
[0.93430825, 0.06569175],
[0.88254513, 0.11745487],
[0.04432996, 0.95567004],
[0.90592932, 0.09407068],
[0.45372052, 0.54627948],
[0.79575333, 0.20424667],
[0.60210556, 0.39789444],
[0.25378242, 0.74621758],
[0.8826708 , 0.1173292],
[0.68806419, 0.31193581],
[0.88265186, 0.11734814],
[0.70004488, 0.29995512],
[0.20716191, 0.79283809],
[0.88806253, 0.11193747],
[0.27558421, 0.72441579],
[0.35745578, 0.64254422],
[0.76086619, 0.23913381],
[0.76666556, 0.23333444],
[0.89997915, 0.10002085],
[0.85337984, 0.14662016],
[0.38572465, 0.61427535],
[0.19051619, 0.80948381],
[0.70391693, 0.29608307],
[0.19484729, 0.80515271],
[0.27212605, 0.72787395],
[0.8839621 , 0.1160379],
[0.74459357, 0.25540643],
[0.23518855, 0.76481145],
[0.81501916, 0.18498084],
[0.38572465, 0.61427535],
[0.90536516, 0.09463484],
[0.4106211 , 0.5893789],
[0.9242273 , 0.0757727],
[0.92791828, 0.07208172],
[0.3104241 , 0.6895759],
[0.40193776, 0.59806224],
[0.68895298, 0.31104702],
[0.91466345, 0.08533655],
[0.07147442, 0.92852558],
[0.87694862, 0.12305138],

[0.38572465, 0.61427535],
[0.6518783 , 0.3481217],
[0.74004035, 0.25995965],
[0.12833996, 0.87166004],
[0.65742258, 0.34257742],
[0.12726699, 0.87273301],
[0.59804896, 0.40195104],
[0.09822878, 0.90177122],
[0.1860191 , 0.8139809],
[0.31348936, 0.68651064],
[0.91906142, 0.08093858],
[0.94989269, 0.05010731],
[0.26444959, 0.73555041],
[0.88593971, 0.11406029],
[0.96157317, 0.03842683],
[0.89670547, 0.10329453],
[0.0781294 , 0.9218706],
[0.65254715, 0.34745285],
[0.91617925, 0.08382075],
[0.96329379, 0.03670621],
[0.80610433, 0.19389567],
[0.94661103, 0.05338897],
[0.44038964, 0.55961036],
[0.18662757, 0.81337243],
[0.34643517, 0.65356483],
[0.9174091 , 0.0825909],
[0.19775673, 0.80224327],
[0.89997915, 0.10002085],
[0.71882255, 0.28117745],
[0.94442549, 0.05557451],
[0.27787712, 0.72212288],
[0.92019239, 0.07980761],
[0.28919633, 0.71080367],
[0.50825781, 0.49174219],
[0.09986118, 0.90013882],
[0.91199107, 0.08800893],
[0.79035999, 0.20964001],
[0.66815148, 0.33184852],
[0.83504206, 0.16495794],
[0.95169134, 0.04830866],
[0.8895021 , 0.1104979],
[0.72377226, 0.27622774],
[0.70821257, 0.29178743],
[0.03599589, 0.96400411],
[0.70495827, 0.29504173],
[0.5906016 , 0.4093984],
[0.73428404, 0.26571596],


```
[0.43514057, 0.56485943],
[0.93660846, 0.06339154],
[0.92241535, 0.07758465],
[0.1338508 , 0.8661492 ],
[0.10301272, 0.89698728],
[0.95516284, 0.04483716],
[0.81402474, 0.18597526],
[0.87695094, 0.12304906],
[0.38541183, 0.61458817],
[0.18340836, 0.81659164]])
```

```
[146]: from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score
```

```
[147]: print("Accuracy: ", accuracy_score(y_test, y_pred))
print("Recall: ", recall_score(y_test, y_pred))
print("Precision: ", precision_score(y_test, y_pred))
print("F1Score: ", f1_score(y_test, y_pred))
```

```
Accuracy:  0.7847533632286996
Recall:    0.6666666666666666
Precision:  0.8
F1Score:   0.7272727272727272
```

```
[148]: # tuning probability cutoff
prob = pd.DataFrame()
```

```
[150]: prob["y_actual"] = y_train
```

```
[152]: prob['p(y=1|x)'] = lr.predict_proba(x_train)[: ,1]
```

```
[153]: prob.head()
```

```
[153]:
```

	y_actual	p(y=1 x)
225	0	0.111281
856	1	0.867430
620	0	0.108423
450	0	0.137570
423	0	0.481504

```
[155]: cut = [float(x)/10 for x in range(0,11)]
cut
```

```
[155]: [0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0]
```

```
[156]: for i in cut:
    prob[i] = prob['p(y=1|x)'].map(lambda x: 1 if x>i else 0)
```

```
[157]: prob.head()
```

```
[157]:
```

	y_actual	p(y=1 x)	0.0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
225	0	0.111281	1	1	0	0	0	0	0	0	0	0	0
856	1	0.867430	1	1	1	1	1	1	1	1	1	0	0
620	0	0.108423	1	1	0	0	0	0	0	0	0	0	0
450	0	0.137570	1	1	0	0	0	0	0	0	0	0	0
423	0	0.481504	1	1	1	1	1	0	0	0	0	0	0

```
[158]: cutoff_df = pd.DataFrame(columns = ['prob', 'accuracy', 'recall', 'precision'])

for i in cut:
    a = accuracy_score(prob['y_actual'], prob[i])
    r = recall_score(prob['y_actual'], prob[i])
    p = precision_score(prob['y_actual'], prob[i])

    cutoff_df.loc[i] = [i,a,r,p]
```


```
[159]: cutoff_df
```

```
[159]:
```

	prob	accuracy	recall	precision
0.0	0.0	0.368263	1.000000	0.368263
0.1	0.1	0.553892	0.943089	0.449612
0.2	0.2	0.715569	0.878049	0.574468
0.3	0.3	0.769461	0.808943	0.650327
0.4	0.4	0.790419	0.731707	0.708661
0.5	0.5	0.821856	0.691057	0.798122
0.6	0.6	0.820359	0.621951	0.850000
0.7	0.7	0.805389	0.512195	0.926471
0.8	0.8	0.773952	0.402439	0.961165
0.9	0.9	0.699102	0.191057	0.959184
1.0	1.0	0.631737	0.000000	0.000000

```
[160]: cutoff_df.plot.line(x = 'prob', y = ['accuracy', 'recall', 'precision'])
```

```
[160]: <AxesSubplot:xlabel='prob'>
```



```
[161]: y_test_pred_prob = lr.predict_proba(x_test)[:,-1]
```

```
[168]: y_test_01 = list(map(lambda x:1 if x>0.45 else 0, y_test_pred_prob))
```

```
[169]: print("Accuracy: ", accuracy_score(y_test, y_test_01))
print("Recall: ", recall_score(y_test, y_test_01))
print("Precision: ", precision_score(y_test, y_test_01))
print("F1Score: ", f1_score(y_test, y_test_01))
```

```
Accuracy:  0.7892376681614349
Recall:    0.6770833333333334
Precision: 0.8024691358024691
F1Score:   0.7344632768361582
```